

Accelerating DNA Read Mapping and Alignment

Fall 2015 - Spring 2016

Sunny Nahar

May 10, 2016

Contents

1	Introduction	6
1.1	What is DNA Sequencing?	6
1.2	Sequencing Platform	6
1.3	Complexity of Read Mapping	6
1.4	Handling Errors	7
1.5	Prefix Trie Method	7
1.6	Seed and Extend Method	7
1.7	Measuring Read Mapping	8
2	Research Question	9
2.1	Improving Seed Selection	9
2.2	Parallelizing Seed Selection	9
2.3	Machine Specs	9
3	Related Work	10
3.1	Uniform Selection	10
3.2	Cheap Kmer Selection (CKS)	10
3.3	Threshold	10
3.4	Hobbes	10
3.5	Optimal Seeds (OSS)	10
3.6	Goal	10
4	Contributions	11
5	Repository structure	12
5.1	Link	12
5.2	Directory	12
5.3	database/inc	12
5.3.1	debug.h	12
5.3.2	functionTimer.h	12
5.3.3	logging.h	12
5.4	database/core	12
5.5	database/seedselection	12
5.5.1	database/seedselection/predictor	13
5.5.2	database/seedselection/selector	13
5.5.3	database/seedselection/test	13
5.5.4	database/seedselection/util	13
5.6	database/util	13
6	Improving the HashTree Data Structure	14
6.1	Introduction	14
6.2	High Throughput Load and Store	14
6.2.1	Design Flaw	14
6.2.2	Working Solution	14

6.2.3	Further Improvements	14
6.3	Memory Allocation	14
6.3.1	Segmented Memory Allocator	15
7	Parallelizing the HashTree	16
7.1	OpenMP Integration	16
7.2	Multithreaded Store	16
7.3	Multithreaded Load	16
7.4	Memory Mapped IO	16
7.5	Greedy Scheduling and Dynamic Work Assignment	17
7.6	Speedup	17
7.7	Future Work	18
7.7.1	Minimizing Disk Thrashing through Reader Thread	18
8	Efficient Space Representations	19
8.1	HashTree	19
8.2	Compressed HashTree	19
8.2.1	Speedup	19
8.3	Compressed HashTree (C structs)	19
8.3.1	Speedup	19
8.4	Future Work	19
8.4.1	Bit level compression	19
8.4.2	Reducing pointer space	20
8.4.3	Succinct and Compact data structures	20
9	Bidirectional Frequency Predictor	21
9.1	Motivation	21
9.2	Description	21
9.3	Implementation	21
9.4	Frequency Average Test	22
9.4.1	Mechanics	22
9.4.2	Selected Results	22
9.4.3	Analysis	23
9.5	Frequency Ratio Test	24
9.5.1	Mechanics	24
9.5.2	Selected Results	24
9.5.3	Analysis	25
9.6	Frequency Comparative Test	26
9.6.1	Mechanics	26
9.6.2	Selected Results	26
9.6.3	Analysis	27
10	Parallel Build of the Frequency Predictor	28
10.1	Building across the HashTree	28
10.2	Synchronization	28
10.3	Speedup	28

11 Seed Selectors	29
11.1 Baseline Methods	29
11.2 Uniform Seed Selector	29
11.2.1 Mechanics	29
11.2.2 Results	29
11.2.3 Analysis	29
11.3 Threshold Seed Selector	29
11.3.1 Mechanics	29
11.3.2 Results	30
11.3.3 Analysis	30
11.4 Hobbes Seed Selector	30
11.4.1 Mechanics	30
11.4.2 Results	31
11.4.3 Analysis	31
11.5 Bidirectional + Hobbes Seed Selector	31
11.5.1 Bidirectional Predictive Seed Extension	31
11.5.2 Mechanics	32
11.5.3 Results	32
11.5.4 Analysis	32
11.6 Bidirectional + Hobbes (Centralized prediction)	32
11.6.1 Centralized Prediction	32
11.6.2 Mechanics	33
11.6.3 Results	33
11.6.4 Analysis	33
11.7 Bidirectional + Uniform (Centralized prediction)	34
11.7.1 Uniform Centralized Prediction	34
11.7.2 Mechanics	34
11.7.3 Results	34
11.7.4 Analysis	34
12 Parallel Seed Selection	36
12.1 Parallelization	36
12.2 Speedup	36
13 Conclusions	37
13.1 Reducing Seed Frequency with the Frequency Predictor	37
13.2 Parallelizing the Pipeline	37
14 Future Work	38
14.1 Iterative Prediction	38
14.2 Optimizing for Comparative Prediction	38
14.3 Extension Patterns	38
14.4 Randomized Division	38
14.5 Probabilistic Models	38
14.6 Machine Learning	39

15 Appendix A	41
15.1 Results of Frequency Average Test	41
15.2 Results of Frequency Ratio Test	47
15.3 Results of Frequency Comparison Test	53
16 Appendix B	56
16.1 Seed Selection Algorithm Speedups	56

1 Introduction

1.1 What is DNA Sequencing?

In the better part of the last decade, DNA sequencing has arisen as an integral tool in the study of biology and medicine. DNA sequencing allows researchers to determine the exact sequence of nucleotides (base-pairs), *Adenine*, *Guanine*, *Thymine*, and *Cytosine*, in a strand of DNA and subsequently determine the sequence of any structure from a small gene to an entire genome. These sequences have a myriad of applications in numerous fields: discovering evolutionary patterns among species, tracing human evolution, identifying genetic components to determine the causes of diseases and construct remedying medicine, and identifying primary characteristics of individuals in forensic studies [1, 2, 3, 4]. This has led to a huge demand for fast and efficient genome sequencing algorithms. Since the technology is becoming a ubiquitous research tool among researchers, universities, and corporations alike, sequencing solutions which do not rely on excessively expensive or custom infrastructure are needed. There is a strong necessity for the ability to sequence DNA on commodity machines and not in days or weeks, but a few hours.

1.2 Sequencing Platform

The advent of largely parallel and high-throughput sequencing (HTS) platforms such as Illumina and SOLiD have revolutionized the field. These platforms sequence DNA by splitting long strands into small chunks of contiguous base-pairs, typically less than 100 base-pairs (bp), called **reads**. Reads are generated in the billions at a very high rate, and the objective is to reconstruct the original DNA sequence by aligning these short reads together [1, 2]. This is an unparalleled challenge on its own, since the computational complexity of aligning billions of reads is enormous. Sequencing the first human genome took more than a decade and over a billion dollars [3]. Fortunately, there are assembled genomes available today, and given the high rate of similarity between genomes, one can be used as a basis for constructing another. Therefore, reads are mapped to a reference genome to construct the overall DNA sequence.

1.3 Complexity of Read Mapping

However, this still poses a very computationally challenging problem. The reference genome is extremely large: the size of the human genome is **3.2 billion** base-pairs long. Reads are short and occur very frequently in the genome, so there are a vast number of possibilities where a single read could map to. In addition, all DNA sequences are not same: different organisms have different DNA structure, and within the same organism, like humans, there can be variations in the DNA due to mutations. This is compounded by possible errors introduced by the sequencing process: insertions, deletions, and substitutions in the read [3]. Therefore, the process of mapping reads to the reference genome is computationally hard since the mapping process needs to account for DNA variation and sequencing errors and allow a certain number of mismatches. The complexity increases with the number of mapping locations and error rate.

Given that there are a multitude of possible mapping locations for a given read, there are two different types of mapping algorithms depending on how many mappings are required. In the first case, a mapper may only need to find some of possible mappings for a read. This is the case in single nucleotide polymorphism (SNP) studies [3], where a few mappings are sufficient. In the other case, all possible mappings for a read may be required. For example in RNA-seq applications, finding all mapping positions is necessary for quantifying the level of gene expression, and in ChIP-seq applications, finding all mapping positions is necessary for characterizing protein binding sites. Mappers are often optimized to handle one of the two approaches [3].

1.4 Handling Errors

There are also multiple standards used for handling read mapping with variation and errors in DNA. Typically, the number of errors in a read is approximately 5% of its length [4]. Therefore there needs to be efficient methods to verify whether a read matches a substring of the reference genome within a specified number of mismatches. Candidate mapping locations for a read are valid only if they pass a similarity metric measuring the number of errors. The first commonly used metric for computing similarity between two strings is **Hamming distance**, the number of differences between one string compared to another. This works well with substitutions, but is not able to capture insertions and deletions effectively. The second metric is **edit-distance**, the smallest number of insertions, deletions, and substitutions needed to change one string to another. This metric is able to capture the full scope of errors in a read. However, this is at the cost of speed, as edit distance calculations are far more complex and expensive than Hamming distance calculations. Algorithms commonly used to compute edit-distance include Smith-Waterman and Needleman-Wunsch [3, 4].

1.5 Prefix Trie Method

There are two primary methods currently used to achieve efficient read mapping: using **suffix arrays/prefix tries** or using **gram-based/seed-and-extend** methods with **hash tables** [3, 1, 2]. Suffix array methods create an efficiently searchable tree data structure consisting of suffixes of the target genome. The tree stores suffixes by encoding characters as edges. A suffix is the concatenation of characters on the path from a leaf to the root. Each node in the tree stores all locations in the reference where that suffix appears. Therefore, searching for mapping locations of a read is equivalent to a search in a tree. In addition, the tree stores information obtained from applying the Burrows-Wheeler Transform and FM-indexing to the reference, allowing the tree to be used as a suffix array. This allows searching the tree with a significantly lower memory footprint. The industrial standard of mappers using this technique include BWA, Bowtie, and SOAP2 [2, 3].

1.6 Seed and Extend Method

In addition to suffix arrays, the other most popular method used is the seed-and-extend / gram-based method. When a mapper is matching a read to the target genome, it must account for errors in the read. Due to the possibility of errors, it is unlikely that the read will match somewhere in the reference exactly. However, it is extremely likely for a set of

substrings of a read to match the genome exactly. These substrings are called **grams** or **seeds**. The mapper chooses a certain number of seeds from a read and searches the genome exactly for these seeds. Once the locations are determined, the mapper tries to determine whether the original read fits at a particular location by extending the seeds into the original read and checking whether it meets the error threshold. To make the algorithms efficient, hash tables are often used to store all possible seeds and their mapping locations in the reference genome. Seeds are often short, typically between 10 – 13 base-pairs in length. Mappers which use this method include MAQ, mrFAST/mrsFAST, SHRiMP, Hobbes, and RazerS [3].

1.7 Measuring Read Mapping

To compare the effective of different algorithms, DNA read mapping methods are generally compared using three metrics. The first is **speed**: how long a mapper takes to map and align a set of reads to the target genome. The second is **sensitivity**: the percentage of reads for which the mapper can find at least one mapping. The last is **comprehensiveness**: the number of correct locations the mapper finds for a given read [3]. Typically, suffix array based mappers are very fast, but have poor sensitivity and comprehensiveness. They are efficient in finding a few mappings as opposed to all possible mappings. Seed-and-extend based methods are typically much slower, but have high sensitivity and comprehensive. They work better at finding all possible mappings of a read.

2 Research Question

Mapper speed remains a bottleneck in DNA read mapping. Processing and aligning billions of reads take a substantial amount of time, and there is a quest to build faster and better algorithms to improve runtime and throughput. In order to further the use DNA sequencing technology, it's important to make it accessible and efficient.

In this research, we focus on improving the complexity and runtime of seed-and-extend based mappers. We focus on two approaches: *developing better seed selection strategies and parallelizing the seed selection infrastructure*.

2.1 Improving Seed Selection

As discussed in section 1.6, seed-and-extend mappers align reads by splitting reads into smaller seeds. These seeds are searched in the reference and used to align the reads. The complexity of the alignment is in correlation with the how often the seeds occur in the reference, or the **frequency of the seeds**.

If seeds occur more frequently, a larger number of locations need to be checked for potential mapping sites of the read, which increases complexity. Therefore it is imperative that we choose seeds which have the **lowest sum of frequencies** to minimize the number of mapping locations, and hence computation.

Therefore the goal is to develop state-of-the-art seed selection schemes which are able to select seeds from a read in a memory efficient and cache efficient manner while minimizing the sum of the frequency of the seeds. These new heuristics and algorithms will be compared with existing seed selection schemes found in seed-and-extend mappers such as Hobbes and FastHash, among others.

2.2 Parallelizing Seed Selection

Single processor compute capabilities have not improved significantly, and with the advent of multi-core and multi-threaded machines, **parallelism** is important in achieving the next generation of performance. Many aspects of DNA read mapping are inherently parallel, so this makes it a viable approach for achieving improvements in runtime.

This research also focuses on optimizing and parallelizing a component of DNA read mapping pipeline, seed selection. We want to determine how efficiently we can distribute work over cores to achieve a large speedup.

2.3 Machine Specs

The machine used for testing is a quad-socket 40 core machine with Intel Hyperthreading and 1000GB of RAM. It has a 64KB L1, 256KB L2, and 24MB L3.

3 Related Work

We now review current seed selection methods.

3.1 Uniform Selection

This is the simplest of selection schemes. The read is split equally into seeds. Therefore this requires no overhead or additional data. Since it does not use any information about seed frequencies or the HashTree it suffers in practice and does not perform well.

3.2 Cheap Kmer Selection (CKS)

This method is used in the FastHASH mapper. The read is split into N contiguous chunks, and the frequency of each chunk is measured through the HashTree. Then the least frequent set of k chunks is selected for k seeds. The seeds can be expanded to fill up the read, since increasing the size of the read does not increase the frequency. By considering k out of the N contiguous chunks, the method is fast, but does not provide much frequency reduction compared to Uniform selection.

3.3 Threshold

This is another common selection method. Given a start position, the length of the seed continually increases until the frequency of the seed, obtainable through the HashTree class, reaches below a given threshold. Seeds are selected contiguously. This method suffers from finding suitable values for this threshold to work well over all reads.

3.4 Hobbes

Hobbes expands on the idea of CKS. It considers all n -length contiguous substrings of the read, and uses dynamic programming to find the optimal set of non-overlapping seeds which are each of length n . This makes it slower as it needs to look up more seed frequencies in the reference and requires a quadratic time dynamic programming calculation. Hobbes performs better than all the previous algorithms in finding a set of seeds with low overall frequency. But it is still limited as it considers only fixed length seeds.

3.5 Optimal Seeds (OSS)

The Optimal Seeds Solver expands on Hobbes by considering variable length seeds. It finds the theoretically optimal division into seeds with the lowest possible frequency. But this is at the price of a high number of HashTree lookups and increased runtime complexity. OSS runs a dynamic programming method to find the optimal divisions as well.

3.6 Goal

Therefore, the goal is to create a hybrid seed selection scheme which performs better than Hobbes, CKS, and Threshold but much faster than OSS, as speed is the limiting factor for OSS.

4 Contributions

This research makes the following contributions:

- Implementation of a general parallel framework for running and testing seed selection algorithms.
- Optimizations on loading, storing, and constructing the HashTree, a substring frequency data structure.
- Implementation and discussion of the Bidirectional Frequency Predictor, a tool to reduce HashTree accesses and speed up seed selection.
- Extensive analysis of the accuracy and viability of frequency predictor.
- A fast method for parallel construction of the frequency predictor.
- Comparative study of the current seed selection schemes.
- Implementation and discussion for several new candidate selection schemes, based on the frequency predictor.

5 Repository structure

5.1 Link

The repository for this project is hosted on GitHub at <https://github.com/xhongyi/buildingblock>.

5.2 Directory

All the necessary code is located in the `database` folder.

5.3 `database/inc`

This includes common utility header files used in most programs.

5.3.1 `debug.h`

This contains assert constructs for style and correctness when debugging. It is enabled through the `-DDEBUGGING` compile time flag.

5.3.2 `functionTimer.h`

This header defines useful functions for measuring time for measuring execution runtime. There is a function for measuring the current wall time and a function for measuring the current clock time. Clock time is aggregated over all threads, so multithreaded execution time should be measured through wall time for physical time and through clock time for compute time.

5.3.3 `logging.h`

Defines logging constructs to be used for execution and debugging. Logging is color coded onto the console for 4 different modes, `INFO`, `WARNING`, `ERROR`, and `DEBUG`. The debug logging is enabled through the `-DDEBUGGING` compile time flag. Enabling the debug logging also adds the file name and function from which the logging is called.

5.4 `database/core`

The folder contains the code for constructing the HashTree representation of the genome. The genome is stored by storing all 30-character long prefixes into a tree, i.e. a prefix trie. The first 10 characters are used to hash into a hashtable which then points to a tree on the remaining 20 characters.

5.5 `database/seedselection`

This folder, which is the main contribution and substance of this research, contains four subfolders, `predictor`, `selector`, `test`, and `util`.

5.5.1 database/seedselection/predictor

This folder contains the implementation of the BiDirectional Frequency Predictor, which is discussed in section 10.

5.5.2 database/seedselection/selector

This folder contains implementations of numerous seed selectors. Further discussion in section 11.

5.5.3 database/seedselection/test

This folder contains tests for the frequency predictor and the seed selectors.

5.5.4 database/seedselection/util

Contains utility classes for generating random reads and seeds and generating seeds based on actual read files such as `ERR240726_1.filt.fastq`.

5.6 database/util

Contains the implementation of a basic memory pool. Discussed in section 6.1.

6 Improving the HashTree Data Structure

6.1 Introduction

The HashTree is a hashtable of suffix tries used to store the genome. It allows for searches with complexity L , where L is the length of the query. More specifically, a search requires L cache misses. To search for a given read in the HashTree, the first 10 base-pairs of the read are used to index into the hashtable, and the remaining base-pairs are used to index into the trie.

6.2 High Throughput Load and Store

A major bottleneck during testing was loading and storing the HashTree representation. Generating the entire tree on the genome was infeasible as it required 4 hours, and this would make testing an arduous process. A export binary format had been created before to speed up this process by writing the HashTree to a file, and reading the HashTree from file during testing. This did speed it up by approximately 2x, to where storing and loading took 2hrs, but it remained a major bottleneck. The binary form of the test genome, `human_g1k_v1.fasta`, was extremely large at 78GB.

6.2.1 Design Flaw

One of the design flaws in the implementation of the load and store functions was small writes to disk interleaved with computation which were creating unnecessary overhead in the write process in the number of function calls.

6.2.2 Working Solution

This was solved by writing the output to the store and load to an internal buffer, which could then be used to read and write directly to the file. This reduced the number of read / write calls to just one.

6.2.3 Further Improvements

Although this provides a performance improvement, this is at the cost of memory usage. Since the entire file is stored in memory, an additional 78GB of memory is required to facilitate that process, whereas the initial method required no overhead. This can likely be reduced by making the buffer the size of the file write buffer, but it depends on the implementation of the file write. Currently, there is not much issue with the memory overhead as the system can support it along with the HashTree.

6.3 Memory Allocation

A critical component of generating the HashTree is allocation of memory. Using the traditional `new` operator causes thread stalls, as `new` contains a lock, and threads have to sequentially request memory. Constructing the HashTree is a memory intensive operation,

with 250GB of memory being requested. Hence contention over memory allocation creates a significant drawback. A custom per thread memory allocator solves this problem. It requests a large chunk of memory using `new`, and facilitates smaller allocations through its own data structure. This allows for contention free memory access.

6.3.1 Segmented Memory Allocator

The solution to the continual use of `new` is creating a type of memory pool, which we call a segmented memory allocator. A memory pool is a small allocator which the program can use instead of `new`. Internally, it requests a large chunk of memory using `new`, and slowly uses up that memory when allocation calls to memory allocator are received.

In this application, the program only needs to allocate and not free memory during execution. This is because the construction of the HashTree only needs to allocate memory for each node but not free. Hence, this simplifies the design of the allocator. As there are no calls to free, memory is not segmented, and hence there are no free blocks to manage or coalesce. Therefore, the memory allocator can simply keep track of a pointer which points to the start of the unused sub-block of memory, and during an allocation call, update the pointer. When the block of memory runs out, it simply calls `new` to allocate another chunk. In addition, freeing memory after execution is much simpler, as only the large chunks need to be free, not the individual requested nodes. So the HashTree keeps a centralized pool of memory requested from all the allocators, which it frees at the end.

Therefore calls to `new` are much faster. This led to a major improvement in execution time of HashTree load, reducing the time from 2 hours to 7 minutes. In addition freeing the memory after using the HashTree becomes much faster since instead of the destructor calling `free`'s on billions of pointers, it only calls `free` on the pointers to each chunk.

The implementation of the allocator is located in `database/util`. The API interface defines a constructor which takes in a chunk size, and an `alloc` function which. Each time the allocator runs out of memory, a new internal block is allocated.

7 Parallelizing the HashTree

7.1 OpenMP Integration

To facilitate parallelization of the interface, the OpenMP framework was used to add multithreading functionality to HashTree operations.

7.2 Multithreaded Store

Parallelizing storing the HashTree as one file is infeasible since the order of bytes in the storage binary matters. Although this could potentially be done by using multiple file descriptors each with its own write location in the file, the number of bytes written by each thread needs to be known beforehand, and this is difficult to obtain correctly.

Therefore the natural extension is to split the HashTree across multiple files, with each thread handling a set of files. This complements well with multithreaded load, where each thread loads from its set of files. As each tree in the hashtable of the HashTree is disjoint, it is natural to split the hashtable among the threads.

Converting each tree into a binary can be done in parallel, since they are independent. However the writes to the file must be such that the entire file is written to before another, to prevent disk thrashing.

We found that using 256 parts for the human genome was the optimal split.

7.3 Multithreaded Load

To load the HashTree, each thread loads its respective file. It is now done through memory mapped IO, discussed in the next section. The files encode multiple tries in the HashTree, and each trie is independent, so each thread can independently construct its respective tries of the HashTree and link it with the others.

Once again, the use of the segmented memory allocator is crucial for performance. Since the memory architecture of the system was nonuniform, a local memory allocator has to be created for each thread, so that cache coherence between processors is not a limiting factor. However being locally allocated, the memory has to be managed as the allocator leaves scope after the thread exits.

7.4 Memory Mapped IO

Reading from disk is inherently slow. Reading roughly 80GB from disk is monumentally worse, especially since it must be read sequentially. When reading from disk, we need each thread to read its file completely. This is to prevent disk thrashing, since if each thread only reads a smaller number of bytes from its file, then the disk will keep seeking between various files. Hence reading is sequential. In addition, it is highly inefficient to operate on data while it is being read. In that case, we have essentially sequentialized the entire process.

So data must be copied from the file to memory before it is used.

This takes significant time and disk bandwidth. A better approach is using Memory Mapped IO. The file is not copied to memory, but mapped to pages in the kernel page cache. Hence we do not need to copy it to user space. Therefore accessing pages is much more efficient. In addition, we can advise the kernel about the page accesses, which in this case is linear. This allows the kernel to prefetch and optimize accordingly. Use memory mapped IO gives a significant boost, as the time to read the file is almost removed.

7.5 Greedy Scheduling and Dynamic Work Assignment

Currently the HashTree is split into 256 parts. There is large variation in file and hence trie sizes, so it is important to distribute work equally among threads. This is done through OpenMP’s dynamic scheduling and a initial greedy distribution of work. OpenMP creates a thread scheduler which dynamically allocates work to available threads: whenever a thread is done with its task, the scheduler gives it the next available piece of work to do. However a problem arises if the sizes of the work in the list is not distributed evenly. For example, if the last piece of work requires a large compute time, then this will likely be sequentially running at the end, as other threads will not have any work to do.

To combat this, we first estimate the size of the work need to load each trie into memory. This is computed by using the file size. We sort the work list of tries from largest to smallest. Therefore the initial batch of work the threads receive are the largest pieces. This is within a factor of 2 to the optimal schedule of work to threads, per the greedy work scheduling algorithm.

7.6 Speedup

The performance caps off at 16 threads, where the speedup is fairly close to linear. This is likely due to the overhead of TLB misses when multiple threads are memory mapping files.

Threads	Runtime (s)	Speedup
1	479.09	1.00
2	232.33	2.06
4	118.89	4.03
8	62.16	7.71
16	35.54	13.48
32	32.23	14.86
64	31.9	15.02
80	31.6	15.16

Table 1: Speedup of building the HashTree across 80 threads

7.7 Future Work

7.7.1 Minimizing Disk Thrashing through Reader Thread

Since there are multiple threads each reading from a respective file, the files are read in arbitrary order compared to the other stored on disk. Hence there is still some disk overhead due to the random access. A solution is to have one reader thread which performs the read operation and notifies other threads when the read has completed.

8 Efficient Space Representations

8.1 HashTree

The HashTree is implemented using `TreeNode` and `LeafNode` classes. These contain extra information used during creation from the genome, so their sizes are large: 48 bytes for the `TreeNode` and 64 bytes for the `LeafNode`. This implementation has a size of 297GB.

8.2 Compressed HashTree

For search queries, we only need to keep track of frequency. Therefore we can create a compressed tree which only encodes this information. Storing the frequency is only 4 bytes, but storing virtual function pointers and struct padding causes the `TreeNode` size to be 48 bytes and `LeafNode` 16 bytes. The `TreeNode` needs to store 4 pointers to each of the subtrees, so this is 32 out of the 48 bytes. This implementation has a size of 186GB.

8.2.1 Speedup

Threads	Runtime (s)	Improvement over HashTree
80	28.47	11%

Table 2: Compressed HashTree across 80 threads

8.3 Compressed HashTree (C structs)

Using structs, we can reduce the space to 4 bytes for the `LeafNode` and 40 bytes for the `TreeNode`, since we no longer virtual function pointers. This reduces the size to 147GB.

8.3.1 Speedup

Threads	Runtime (s)	Improvement over HashTree
80	26.12	21%

Table 3: Compressed Struct HashTree across 80 threads

8.4 Future Work

8.4.1 Bit level compression

We observe that the frequency requires at most 22 bits. Hence the extra 10 bits in the 4 byte integer used to store the frequency is wasted. This can potentially be used to store additional information to aid with compression. This is detailed below.

8.4.2 Reducing pointer space

Storing 4 8-byte pointers requires a lot of space. We detail a scheme which can reduce this to storing 1 pointer. Given a `TreeNode`, if we store all child nodes successively, then only the start pointer to the first node is required. The other 3 can be indexed using the first pointer. Therefore we need an 8-byte pointer to point to the start. We need to keep track of which children (ACGT) exist, so this requires 4 indicator bits. Hence to calculate the position of the child, we add the corresponding indicator bits and multiply by 8 to get the offset. However, the child node sizes are not uniform, since the `LeafNode` is 4 bytes, but the `TreeNode` is at least 16 bytes (8 bytes for pointer, 4 bytes for frequency and indicator bits, and 4 bytes for padding). So we also need to keep track of whether the child is a `LeafNode` or `TreeNode`. This takes an additional 4 bits. Hence we are able to use 16 bytes (with 4 bytes and 2 bits free) to store the `TreeNode` and 4 bytes for the `LeafNode`.

8.4.3 Succinct and Compact data structures

There has been considerable research in building very compact representations of binary trees which use 2-4 bits per node in addition to the data and still allow for a searchable representation. It is probably possible to reduce the pointer to perhaps an integer or less.

9 Bidirectional Frequency Predictor

9.1 Motivation

The crux of the development of this research is through the concept of the predicting frequency. We want to be able to minimize queries to the HashTree for frequency of seeds and instead create a data structure which essentially provides a dimensionality reduction over the space of seeds.

We can view seeds (up to length 30) as a 31 dimensional vector in a space \mathcal{S} of size $30 * 4^{30}$, where the first dimensional specifies length, and the remaining specify the seed. We have a function $f : \mathcal{S} \mapsto \mathbb{N}$ which maps a vector and its corresponding seed to the frequency of the seed in the reference. Clearly storing every possible seed frequency requires too much space. Instead, we want to store less information such that we can get an estimate of the frequency.

9.2 Description

Therefore, we create a data structure called a bidirectional frequency predictor. This is a table which given a base seed, left extension, and right extension, gives an estimate of the frequency of the seed if the base was extended to the left by a certain number of characters and to the right by a certain number of characters. Therefore if the base seed has length 10, we can support up to a left and right extension of 20, given the max length as 30, so the size of the table is $4^{10} * 20 * 20 * 4 = 1.6\text{GB}$, a much more manageable size. For a table with base seed length 12 and left and right extensions length of 18, the size is 20GB.

Note that it is very difficult for the prediction to be accurate. This is due to many factors. The prediction for a given base seed and left and right extension is made by averaging out all paths in the HashTree which are the desired length and have the base seed in the appropriate place. As the frequencies of seeds are heavily variable, ranging in 4 - 5 orders of magnitude even at the same depth, averages are skewed. Therefore the expectation here is that if the prediction is within a factor of 10x, maybe even 20-50x, it may be useful.

9.3 Implementation

The implementation is located in `database/seedselection/predictor`. The predictor is created by traversing the HashTree and updating the respective part of the table during the traversal. The frequency of a node is weighted by itself during the average calculation. This is because the frequency represents the number of positions in the reference which have that string occurring, and a highly frequent seed is more likely to appear, so it gets higher weight. Once the predictor is created, it can be stored to a file so that it does not have to be recreated each time.

9.4 Frequency Average Test

9.4.1 Mechanics

The first test used the `ERR240726_1.filt.fastq` read set composing of 4026809 reads of length 100. To get k -length seeds, the first k letters of each read were taken. The seeds were divided into 12 sets based on the actual frequency: [0], [1-5], [6 - 10], [11 - 50], [51 - 100], [101 - 500], [501 - 1000], [1001 - 5000], [5001 - 10000], [10001 - 100000], [100001 - 1000000], [1000001 - 1000000000]. For each set, the number of seeds in each set were counted, the average exact frequency of the seeds was calculated. In addition, for each possible sets of left and right extensions, the average predicted frequency was calculated. This was predicted through a bidirectional frequency predictor with base seed length of 10. Therefore, for example, (3, 4) in the tables below indicates that the prediction of a length $10 + 3 + 4 = 17$ seed was calculated such that positions 3 – 12 were used for the base seed and searched in the table with left extension 3, right extension 4.

9.4.2 Selected Results

Full results are available in Appendix A.

Freqs:	0	5	10	50	100	500
Count:	276	7852	15058	214606	155766	2045663
Avg Freq:	0.00	3.40	8.24	28.44	76.68	276.09
(0, 2):	15.93	36.43	60.24	140.76	297.01	564.25
(1, 1):	17.46	39.70	62.21	143.90	284.38	558.51
(2, 0):	16.09	30.70	57.89	145.43	290.71	566.10

Table 4: Seed length 12

Freqs:	1000	5000	10000	100000	1000000	1000000000
Count:	688840	555620	108938	179399	52030	2761
Avg Freq:	693.54	2074.50	7095.21	36617.05	196559.76	2059168.57
(0, 2):	1234.97	3751.62	9835.89	39765.91	189361.26	1467818.82
(1, 1):	1207.22	3685.46	10587.19	38192.25	201147.23	1478942.37
(2, 0):	1242.16	3843.22	10190.89	42736.41	194746.92	1470746.04

Table 5: Seed length 12

Freqs:	0	5	10	50	100	500
Count:	119893	1296668	782215	1012545	169448	228469
Avg Freq:	0.00	2.90	7.71	21.54	70.75	231.81
(0, 5):	94.76	142.80	196.01	378.19	1050.82	2593.59
(1, 4):	88.25	137.32	196.19	365.75	1028.75	2575.29
(2, 3):	92.72	139.16	192.85	357.71	1051.37	2620.70
(3, 2):	93.05	139.42	187.77	364.28	1062.28	2669.95
(4, 1):	97.40	139.39	190.59	373.66	1074.14	2725.39
(5, 0):	92.20	143.02	198.27	391.13	1134.53	2930.94

Table 6: Seed length 15

Freqs:	1000	5000	10000	100000	1000000	1000000000
Count:	67449	147865	68003	110002	22731	1521
Avg Freq:	708.27	2513.11	7125.24	38015.44	173055.18	1191151.41
(0, 5):	6500.26	9933.95	11493.09	33278.47	123185.21	499685.63
(1, 4):	5368.47	9709.77	12161.65	33446.83	131134.11	509271.39
(2, 3):	5750.41	9122.51	12469.85	34774.53	134129.87	514628.16
(3, 2):	5787.67	9534.10	12678.14	34910.12	137436.07	518084.84
(4, 1):	5965.05	10939.14	13353.92	33978.20	142319.18	519248.13
(5, 0):	7576.56	11455.24	13135.02	34583.92	138267.35	513356.68

Table 7: Seed length 15

9.4.3 Analysis

For a particular seed length, we can see that the prediction is best for seeds with higher frequencies and gets progressively worse as the seed frequency decreases. This is due to the average being skewed toward large frequencies. The weights of the seeds with higher frequency is larger than the weights with lower frequency for the same total length, so seeds with lower frequencies will be overestimated. However the assumption is that since these seeds have low frequencies, we do not expect to see too many of them, so the prediction will be mostly accurate. Using a predictor with base length 10, we see that for seeds with length at most 15 with frequency at least 100, the prediction is at most 13x away. For seeds with length at most 30 with frequency at least 100, the prediction is at most 30x away.

9.5 Frequency Ratio Test

9.5.1 Mechanics

In this test, we compute the ratio of the predicted frequency to the exact frequency for seeds in the read set `ERR240726_1.filt.fastq`. The average ratio is displayed over different seed lengths and frequency classes. The classes are the same as the previous test.

9.5.2 Selected Results

Full results are available in Appendix A.

Freqs:	0	1	5	10	50	100
Count:	276	889	6963	15058	214606	155766
Avg Freq:	0.00	1.00	3.70	8.24	28.44	76.68
(0, 2):	0.00	16.04	10.40	7.33	5.07	4.00
(1, 1):	0.00	24.43	11.36	7.63	5.20	3.81
(2, 0):	0.00	18.23	9.17	7.12	5.17	3.95

Table 8: Seed length 12

Freqs:	500	1000	5000	10000	100000	1000000000
Count:	2045663	688840	555620	108938	179399	54791
Avg Freq:	276.09	693.54	2074.50	7095.21	36617.05	290419.39
(0, 2):	2.13	1.78	1.81	1.40	1.16	0.97
(1, 1):	2.13	1.74	1.78	1.50	1.20	1.04
(2, 0):	2.13	1.79	1.84	1.45	1.26	1.00

Table 9: Seed length 12

Freqs:	0	1	5	10	50	100
Count:	119893	287608	1009060	782215	1012545	169448
Avg Freq:	0.00	1.00	3.45	7.71	21.54	70.75
(0, 5):	0.00	104.66	48.59	25.94	18.37	14.84
(1, 4):	0.00	104.44	46.24	25.88	17.93	14.48
(2, 3):	0.00	107.86	46.88	25.49	17.33	14.86
(3, 2):	0.00	106.20	47.02	24.87	17.67	15.08
(4, 1):	0.00	106.30	46.82	25.13	18.16	15.23
(5, 0):	0.00	105.99	48.37	26.19	18.84	16.19

Table 10: Seed length 15

Freqs:	500	1000	5000	10000	100000	1000000000
Count:	228469	67449	147865	68003	110002	24252
Avg Freq:	231.81	708.27	2513.11	7125.24	38015.44	236906.59
(0, 5):	11.64	9.26	4.70	1.64	1.01	0.70
(1, 4):	11.91	7.65	4.48	1.75	1.01	0.75
(2, 3):	12.03	8.19	4.31	1.78	1.06	0.77
(3, 2):	12.20	8.27	4.44	1.81	1.09	0.79
(4, 1):	12.54	8.52	5.05	1.92	1.07	0.82
(5, 0):	13.13	10.86	5.42	1.87	1.09	0.80

Table 11: Seed length 15

9.5.3 Analysis

In this test, we have a clearer picture of by what factor the predicted frequencies differ from the actual frequencies. As this is just a reinterpretation of the previous data, the trends are the same. The prediction gets progressively worse as the frequency in a group of seeds of length m decreases. For seeds with length at most 15, we can predict seeds with large frequency (at least 10000) within a factor of 2. Expanding this to seeds with frequency at least 1000, we can predict this within a factor of 10. For seeds with length at most 20, the predictive ratios are slightly better: within a factor of 1.9 with frequency at least 10000, within a factor of 7.7 for seeds with frequency at least 1000, and within a factor of 30 for seeds with frequency at least 100. For seeds with length at most 25, the ratio is within 4.3 for seeds with frequency at least 1000. So the predictor works fairly well for high frequency seeds.

9.6 Frequency Comparative Test

9.6.1 Mechanics

This test is a comparative test between seeds. The test takes roughly 4 million pairs of seeds and calculates the percentage of the time the predictor predicts the correct sign, i.e. the percentage of time if the actual frequency of seed 1 is larger than seed 2 the predicted frequency of seed 1 is larger than seed 2, and vice versa.

9.6.2 Selected Results

Full results are available in Appendix A.

(left, right)	% of correct comparisons
(0, 2):	0.838487
(1, 1):	0.834117
(2, 0):	0.838308

Table 12: Seed length 12

(left, right)	% of correct comparisons
(0, 5):	0.734238
(1, 4):	0.733823
(2, 3):	0.734016
(3, 2):	0.734292
(4, 1):	0.733973
(5, 0):	0.735255

Table 13: Seed length 15

(left, right)	% of correct comparisons
(0, 8):	0.644817
(1, 7):	0.648471
(2, 6):	0.647196
(3, 5):	0.647320
(4, 4):	0.649201
(5, 3):	0.649550
(6, 2):	0.649888
(7, 1):	0.649108
(8, 0):	0.649469

Table 14: Seed length 18

9.6.3 Analysis

As expected, the accuracy decreases with increasing seed size. This is due to, for larger length seeds, an average taken over a larger and more variable distribution, which will not be very accurate. The accuracy was consistent over the extension combinations for each seed length, differing at most by 0.01%. The predictor was 88% accurate for 11-length seeds, at least 73% accurate for seeds with length at most 15, and at least 60% accurate for seeds with length at most 25%. The predictor's high percentage in comparative accuracy makes it a useful tool in deciding which seeds to pick.

10 Parallel Build of the Frequency Predictor

10.1 Building across the HashTree

The frequency predictor is built by traversing the HashTree and updating a global table at each iteration. Each trie of the HashTree can be traversed in parallel by a thread, so this makes the build operation highly parallel.

10.2 Synchronization

A simple method to update the global table would be to grab a lock or grab a per-element lock. Acquiring a lock can potentially take time and requires memory to store all the lock data structures. A simpler method is using OpenMP's atomic update. This essentially performs a CAS (compare and swap) to update the table.

Since the table is huge and we have at most 80 threads concurrently modifying it, it is unlikely that two threads will access the same element in the same. Therefore the CAS should succeed in almost all cases.

It was observed that removing the atomic write improved the performance by 30% and still produced the correct table output, but this is still unsafe in the general case.

10.3 Speedup

The speedup is linear up to 16 threads, but slowly degrades to about half the performance at 80 threads. The use of atomic writes is a contributing factor, as removing the atomic requirement substantially increases speed, but at the cost of errors. However no errors were observed.

Threads	Runtime (s)	Speedup
1	2930	1.00
2	1470	2.00
4	727	4.03
8	364	8.07
16	190	15.42
32	109	26.88
64	82	35.73
80	71	41.27

Table 15: Speedup of building the Predictor across 80 threads

11 Seed Selectors

11.1 Baseline Methods

Three seed selectors using existing methods were implemented to compare the performance of the predictor: Uniform, Threshold, and Hobbes. Both Uniform and Threshold are used as a starting point for the seed selectors using the predictor, so they also provide a direct comparison in the percentage reduction of seed frequency. The seed selectors were compared in both frequency sum reduction and runtime. Comparison of runtime is subjective since there are always performance improvements which can be made to optimize the seed selection algorithm. Therefore absolute comparison of runtime is only somewhat meaningful. Since the predictive algorithms use Hobbes and Uniform seed selection schemes as a baseline, they can be compared reliably to their non-predictive variant.

11.2 Uniform Seed Selector

11.2.1 Mechanics

The Uniform Seed Selector simply chooses equal length seeds. It does not have any parameters. The Uniform Seed Selector was tested with the `ERR240726_1.filt.fastq` read set with approximate 4 million reads of size 100 and 5 seeds.

11.2.2 Results

Initialization (s)	Runtime (s)	Average sum of frequencies
0.00	0.83	5247.36

Table 16: Uniform Seed Selector results (5 seeds)

11.2.3 Analysis

The average sum of frequencies of the seeds in a read was 5247.36 when they were chosen equally.

11.3 Threshold Seed Selector

11.3.1 Mechanics

The Threshold Seed Selector increases the seed length of a seed until it falls below a certain threshold. It takes in as a parameter the threshold at which to stop extending the seed. The Threshold Seed Selector was tested with the `ERR240726_1.filt.fastq` read set with approximate 4 million reads of size 100. The threshold parameter was varied to within reasonable bounds.

11.3.2 Results

Threshold	Initialization (s)	Runtime (s)	Average sum of frequencies
100	0.00	32.75	7804.55
200	0.00	30.34	6738.40
300	0.00	28.49	6281.90
400	0.00	27.33	6038.36
500	0.00	26.71	5914.05
600	0.00	25.65	5881.85
700	0.00	25.15	5894.80
800	0.00	24.72	5945.74
900	0.00	24.29	6022.96
1000	0.00	23.92	6129.13

Table 17: Threshold Seed Selector results (5 seeds)

11.3.3 Analysis

In all cases, the Threshold seed selector performs worse than the Uniform seed selector. It reaches its optimum at the threshold of 600 and increases on both sides. This is due to the fact that if the threshold is too small, then too many base pairs will be used initially, forcing the last seeds to have the minimal possible length, which have extremely high frequency. If the threshold is too large, valuable base pairs are essentially unused and wasted on the end seeds which do not seed them. Clearly the threshold predictor is not a good choice, especially if it cannot perform better than the naive method. It is also much slower than the Uniform seed selector, by a factor of 30x.

11.4 Hobbes Seed Selector

11.4.1 Mechanics

The Hobbes Seed Selector uses a dynamic programming algorithm to figure out the optimal set of k seeds where all the seeds have a common fixed length. This is a parameter the Hobbes Seed Selector takes in. Since the seeds given by Hobbes do not always use up the entire read, they are extended uniformly to fill up the read. If there were extra base pairs available between two seeds, the seeds were extended equally to use the extra base pairs. The Hobbes Seed Selector was tested with the `ERR240726_1.filt.fastq` read set with approximate 4 million reads of size 100 and 5 seeds with varying values of the base seed length parameter.

11.4.2 Results

Hobbes seed length	Initialization (s)	Runtime (s)	Average sum of frequencies
10	0.00	161.98	1957.68
11	0.00	328.19	1470.05
12	0.00	383.13	1293.36
13	0.00	468.93	1232.43
14	0.00	524.73	1254.66
15	0.00	574.62	1350.23
16	0.00	638.47	1499.83
17	0.00	695.79	1759.31
18	0.00	699.05	2261.20
19	0.00	683.73	3302.05

Table 18: Hobbes Seed Selector results (5 seeds)

11.4.3 Analysis

The frequency reduction of Hobbes increases with Hobbes seed length until seed length 13, at which it reduces at a much higher rate until seed length 19. The runtime increases with the Hobbes seed length. It jumps in the beginning, double at length 10 to 11, but becomes linear after that. Hobbes performs much better than the Uniform and Threshold Seed Selectors. The fastest Hobbes (seed length 10) is about 2.5x better than Uniform and Threshold, although it is 7x slower than Threshold, and considerably slower than Uniform. The best Hobbes selector (seed length 13) performs 4.5x better than Uniform and Threshold, but it is 21x slower. Since the maximum Hobbes seed length is 20, given read length of 100 and 5 seeds, as the seed length becomes close to the maximum, it is likely that many seeds use extra base pairs which they do not need. Then other seeds remain too frequent, causing the increase in frequency we see above.

11.5 Bidirectional + Hobbes Seed Selector

11.5.1 Bidirectional Predictive Seed Extension

The first of the predictive seed selectors is the Bidirectional Hobbes Seed Selector. Using a parameter describing the length of each seed k , it calls Hobbes to find the optimal set of seeds of length k in the read. Then instead of performing uniform seed extension in Hobbes, we perform a predictive seed extension. Hence during extension, no calls to HashTree are made. Suppose there are d base pairs between seeds i and $i+1$. Then there are $d+1$ possible ways to split the base pairs among seeds. For each division (a, b) where $a + b = d + 1$, the frequency is estimated using the predictor, where the Hobbes seed is used as the base seed for the predictor. The right extension for seed i is a , and the left extension for seed $i + 1$ is b . The division with the lowest predicted sum of seeds is picked. This is processes linearly from left to right among the seeds, first optimizing the divider between seed 1 and 2, and then 2 and 3, and so forth.

Note: refer to section 10 for reference on predictor terminology.

11.5.2 Mechanics

The Bidirectional Hobbes Seed Selector was tested with the `ERR240726.1.filt.fastq` read set with approximate 4 million reads of size 100 and 5 seeds. Since the Hobbes seeds are used as anchors and cannot move and the predictor uses the Hobbes seeds as the base seeds, the predictor base length must be the same as the Hobbes length. Since we are only testing with the base length 10 predictor, we can only test with the Hobbes length 10 as of now. We test using the predictor and using the HashTree for exact frequencies in order to provide a measurement of the effectiveness of the predictor.

11.5.3 Results

Hobbes Length	Initialization (s)	Runtime (s)	Average sum of frequencies
10	3.64	234.80	1817.16

Table 19: Bidirectional Hobbes Seed Selector results (5 seeds) (Predictive frequencies)

Hobbes Length	Initialization (s)	Runtime (s)	Average sum of frequencies
10	0.00	319.20	1788.83

Table 20: Bidirectional Hobbes Seed Selector results (5 seeds) (Exact Frequencies)

11.5.4 Analysis

Compared to the Hobbes seeds selector with Hobbes length 10, this has a 7% reduction in frequency sum of the seeds. The increase in runtime is about 45%. Without data for other Hobbes lengths, it is hard to determine how effective this method is over Hobbes. It clearly outperforms Threshold and Uniform as well. The predictive version is only 2% away from the optimal division using exact frequencies, as shown above. This shows that the predictor performs fairly well in choosing seed division close to the optimal for this method. As the theoretically lowest frequency sum using this method with a Hobbes seed length of 10 is only 2% lower, it seems as if this does not provide much potential for frequency reduction.

11.6 Bidirectional + Hobbes (Centralized prediction)

11.6.1 Centralized Prediction

This seed selector is a variant of the previous one. We again use Hobbes to find the initial set of seeds of length k . When we perform seed extension with the remaining base pairs, we use a centralized prediction scheme: if we are looking at the extension of the Hobbes seed by a base pairs to the left and b base pairs to the right, the instead of using the Hobbes seed as the predictor base seed and a as the left extension and b as the right extension, we

choose the middle k of the $a + k + b$ base pairs as the predictor base seed, with left extension $(a + b)/2$ and right extension $(a + b)/2$. There are more possible divisions in this method since the predictor base seed is not restricted to the Hobbes "anchor" seeds as before. Once again, updates to the divisions are performed linearly from left to right.

Note: refer to section 10 for reference on predictor terminology.

11.6.2 Mechanics

The Centralized Bidirectional Hobbes Seed Selector was tested with the `ERR240726_1.filt.fastq` read set with approximate 4 million reads of size 100 and 5 seeds. Hobbes seeds are no longer anchors in this method, we can easily test over different Hobbes seed lengths with the same predictor. The results here are shortened due to large amounts of computational time required for the higher Hobbes seed lengths.

11.6.3 Results

Hobbes seed length	Initialization (s)	Runtime (s)	Average sum of frequencies
10	3.60	305.14	1475.86
11	3.60	447.17	1396.98
12	3.60	524.43	1371.63
13	3.60	588.36	1369.19
14	3.60	644.77	1388.12
15	3.60	703.14	1425.18
16	3.60	735.22	1472.08

Table 21: Centralized Bidirectional Hobbes Seed Selector results (5 seeds) (Predictive Frequencies)

Hobbes seed length	Initialization (s)	Runtime (s)	Average sum of frequencies
10	0.00	859.56	1008.44
11	0.00	1024.30	940.11

Table 22: Centralized Bidirectional Hobbes Seed Selector results (5 seeds) (Exact Frequencies)

11.6.4 Analysis

Centralized Bidirectional Hobbes performs the best out all the predictive seed selectors. Using a Hobbes length of 10, its average frequency sum is 1476, 10% better than Centralized Bidirectional Uniform and 19% better than Bidirectional Hobbes. It is 25% better than traditional Hobbes. Its runtime is about 33% more than Bidirectional Hobbes and 50% more than Centralized Bidirectional Uniform. Although surprisingly, it does not do better than Hobbes for Hobbes length of 12 - 15. This is an interesting result and cause is still undetermined. The theoretically optimal frequency using this method for Hobbes length 10 is 1008.44, which is lower than all of Hobbes. If the predictor is able to be improved, this method would perform really well.

11.7 Bidirectional + Uniform (Centralized prediction)

11.7.1 Uniform Centralized Prediction

In this seed selector, we use the Uniform seed selector as the baseline, so seeds are split equally in the read. Then for each pair of adjacent seeds, we view the conjunction of the base pairs as one large chunk of base pairs, and try all possible divisions of this chunk into two seeds. The optimal would be picked by choosing the one with the lowest predicted sum of frequencies, using centralized prediction, as detailed in the previous seed selector in section 11.6.1.

For example with a read length of 100 and 5 seeds, each seed would initially start out with 20 base pairs. In the first iteration, the sequence of 40 base pairs of the first two seeds would be looked at. If the minimum seed length was 10, then the following divisions would be considered: (10, 30), (11, 29), \dots , (29, 11), (30, 10).

11.7.2 Mechanics

The selector was tested with the `ERR240726_1.filt.fastq` read set with approximate 4 million reads of size 100 and 5 seeds. Results were calculated using both the predicted frequency from the predictor and the exact frequency from the HashTree to compare effective of the predictor and runtime.

11.7.3 Results

Initialization (s)	Runtime (s)	Average sum of frequencies
3.53	205.66	1607.09

Table 23: Centralized Bidirectional Uniform Seed Selector results (5 seeds) (Predictive frequencies)

Initialization (s)	Runtime (s)	Average sum of frequencies
0.00	796.38	1359.63

Table 24: Centralized Bidirectional Uniform Seed Selector results (5 seeds) (Exact frequencies)

11.7.4 Analysis

The predictive seed selector performed much better than its Uniform counterpart, reducing the frequency sum by a factor of 3.3x. This was at the cost of a 200s increase in computation time. It performed better than the Bidirectional Hobbes seed selector, which is surprising, since the expectation was that Hobbes seeds were the optimal k length seeds, so they should be part of the optimal overall seeds. It was about 11% better than Bidirectional Hobbes. It also ran 14% faster, as no Hobbes calculations needed to be done. Comparing this to the version with exact frequencies, the theoretically lowest possible frequency sum was 15.6%

lower. So in this case, using the predictor was not as accurate as in Bidirectional Hobbes, but reasonable accurate. Since there is a gap, it is likely that we can reduce the predictive version lower with improvements to the predictor and reach the lower bound.

12 Parallel Seed Selection

12.1 Parallelization

This problem is almost embarrassingly parallel. We can simply parallelize over reads in the set, as there are millions of such reads, instead of trying to parallelize the selection algorithms themselves, which do not exhibit much parallelism.

The seed selection algorithms access the HashTree and Predictor concurrently. Since these are static data structures, multiple threads can concurrently read the data. However there will be overhead due to cache coherence traffic from the reads propagating between caches which will bring down the overall performance.

In addition, the HashTree is large, so it is spread across multiple NUMA nodes. This implies that there will be effects due to nonuniform memory accesses. Threads located closer to the core whose memory node contains the HashTree will perform much faster than threads located faraway. So we will definitely expect to see an impact due to NUMA on the performance. The frequency predictor is likely stored across 1 node, so threads on other nodes have higher memory latencies.

12.2 Speedup

Almost every selector was able to get high speedup across the multicore system. The Uniform Seed Selector performed poorly, mainly due to extremely low arithmetic intensity, so the overhead of multicore multithreading exceeded the benefits. The Centered Bidirectional Uniform Selector, which only accessed the frequency predictor, had a 43x speedup over 80 threads. The remaining seed selectors which accessed both the predictor and the HashTree had speedups between 60x - 67x over 80 threads, which is good given the numerous limitations, most notably the NUMA architecture. Full results are located in Appendix B.




Figure 1: Speedup of seed selectors over 80 threads

13 Conclusions

13.1 Reducing Seed Frequency with the Frequency Predictor

The Bidirectional Frequency predictor was able to reduce the frequency sum of seeds in most cases. We investigated predictor specific tests in section 9, and found that the predictor did well in predicting seeds with high frequencies. Ultimately, we do not want there to be any high frequency seeds when the read is split into seeds. It is much better for the predictor to perform well on high frequency seeds and not as well on low frequency seeds since high frequency seeds are several magnitudes more expensive. Also the predictor functioned well in the comparative sense and was able to judge whether a seed had higher frequency than another with 60 – 88% accuracy, depending on the seed length.

When we added the Bidirectional predictor to Hobbes and the Uniform seed selectors, they showed marked improvements in most cases. The predictor was able to achieve a significant reduction in the sum of the frequencies without much overhead in runtime and initialization time.

There are a considerable number of possible improvements to the predictor which could increase its performance and accuracy. This method has shown to be a viable candidate in lowering the frequency sum of seeds. It has a lot of potential for future work.

13.2 Parallelizing the Pipeline

We were able to improve the runtime and effectively parallelize the entire Seed Selection pipeline, from building the HashTree, constructing the predictor, and performing Seed Selection using various algorithms.

- There was a 15x improvement to loading the HashTree from disk. We described potential ways to decrease space and improve the runtime, further lowering the overhead of the HashTree load.
- There was about a 42x improvement to constructing the Bidirectional Frequency predictor. We can build larger predictors much more efficiently and even on the fly, since the small build time likely offsets the cost of storing on disk.
- We achieved a very high efficiency in parallelizing seed selection, averaging about a 60x speedup over 80 threads. There is room for further improvement if we analyze and optimize for NUMA accesses using thread affinity.
- Reducing the runtime for seed selection speeds the overall mapping process, and allows for more complex, expensive, but accurate seed selection algorithms for the same cost as running basic algorithms sequentially.

14 Future Work

There are a wide variety of future avenues to consider in this research. Developing the frequency predictor was only the first step in determining the effectiveness of this method. This first iteration showed that there is some promise to this method if it can be improved and optimized.

14.1 Iterative Prediction

In each of the predictive seed selectors, we run a linear pass of the seed division scheme which optimizes the divisions of two adjacent seeds. This is only done once. It is plausible that if this is repeated multiple times, the frequency would reduce further as the optimal divisions would get updated. However the reduction may not be by much, and this would not warrant the extra time needed to compute the divisions. Therefore, it could be useful to analyze how multiple iterations affect the frequency reduction.

14.2 Optimizing for Comparative Prediction

The earlier results showed that the predictor worked reasonably well in determining if one seed had a larger frequency than other. This was not originally the purpose of the predictor, and it could be possible to build a predictor specifically for this purpose and determine the performance benefits of using a comparative prediction versus a numerical prediction.

14.3 Extension Patterns

By examining the predictive frequency of a seed across the various extensions, there is a pattern that emerges. For seeds which have large variation in the predicted frequency, the variation is not randomly distributed, but appears in chunks. This is because the chunk corresponds to a certain substring of the seed which has high frequency. By analyzing the different frequencies across the extensions, it is likely that we can create an much more accurate aggregated guess instead of directly using the predictor.

14.4 Randomized Division

In the implementations of the predictors, the seed extension was performed from left to right, so that the division of the second seed was dependent on the first, and the division of the third seed was dependent on the first two, and so forth. This creates a bias in the range of possible divisions. A possible solution is randomizing which seed division we optimize.

14.5 Probabilistic Models

Considering second order statistic like standard deviation may be helpful instead of just considering the average. By using the mean and standard deviation and assuming a rough distribution, we would be able to construct a probabilistic distribution of the possible frequency of a read, and combine these distributions together to form the probabilistic model for the frequency of all the seeds of a read.

14.6 Machine Learning

Ultimately, the problem here which the predictor is solving is predicting the frequency of substrings in a long string. As mentioned before, we want to learn a model to estimate the function which maps seeds to frequency. It is possible that there may be structure in the seed frequencies, and they could be predicted from a machine learning model using a certain set of features. This is an interesting area which could show promise.

References

- [1] Ahmadi A, Behm A, Honnalli N, Li C, Weng L, and Xie X. Hobbes: optimized gram-based methods for efficient read alignment. *Nucleic Acids Research*, 40(6), 2012.
- [2] Heng Li and Richard Durbin. Fast and accurate short read alignment with burrows-wheeler transform. *BioInformatics*, 25(14):1754–1760, 2009.
- [3] Hongyi Xin, Donghyuk Lee, Farhad Hormozdiari, Samihan Yedkar, Onur Mutlu, and Can Alkan. Accelerating read mapping with fasthash. *BMC Genomics*, 14.
- [4] Hongyi Xin, Sunny Nahar, Richard Zhu, John Emmons, Gennady Pekhimenko, Carl Kingsford, Can Alkan, and Onur Mutlu. Optimal seed solver: Optimizing seed selection in read mapping. *BioInformatics*.

15 Appendix A

15.1 Results of Frequency Average Test

Freqs:	0	5	10	50	100	500
Count:	3	370	1374	25497	79449	396593
Avg Freq:	0.00	3.83	8.21	34.83	76.60	296.36
(0, 1):	7.00	20.69	27.90	87.38	174.12	525.32
(1, 0):	7.33	18.72	27.47	87.28	174.09	528.46

Table 25: Seed length 11

Freqs:	1000	5000	10000	100000	1000000	1000000000
Count:	924687	1989959	261184	273495	70693	3505
Avg Freq:	758.52	2063.03	6921.88	32482.25	204641.16	2457548.71
(0, 1):	1038.68	2524.14	8008.69	32842.72	208485.70	2074929.54
(1, 0):	1039.14	2542.68	8089.95	32705.69	216075.70	2074192.63

Table 26: Seed length 11

Freqs:	0	5	10	50	100	500
Count:	276	7852	15058	214606	155766	2045663
Avg Freq:	0.00	3.40	8.24	28.44	76.68	276.09
(0, 2):	15.93	36.43	60.24	140.76	297.01	564.25
(1, 1):	17.46	39.70	62.21	143.90	284.38	558.51
(2, 0):	16.09	30.70	57.89	145.43	290.71	566.10

Table 27: Seed length 12

Freqs:	1000	5000	10000	100000	1000000	1000000000
Count:	688840	555620	108938	179399	52030	2761
Avg Freq:	693.54	2074.50	7095.21	36617.05	196559.76	2059168.57
(0, 2):	1234.97	3751.62	9835.89	39765.91	189361.26	1467818.82
(1, 1):	1207.22	3685.46	10587.19	38192.25	201147.23	1478942.37
(2, 0):	1242.16	3843.22	10190.89	42736.41	194746.92	1470746.04

Table 28: Seed length 12

Freqs:	0	5	10	50	100	500
Count:	3802	98737	98153	749457	1001262	1375999
Avg Freq:	0.00	3.16	7.87	33.10	73.31	206.49
(0, 3):	28.29	66.67	128.05	205.02	310.86	686.97
(1, 2):	33.65	64.81	109.04	206.42	309.03	669.90
(2, 1):	43.19	61.63	118.76	210.73	315.53	668.94
(3, 0):	32.25	63.77	129.10	207.58	313.69	700.56

Table 29: Seed length 13

Freqs:	1000	5000	10000	100000	1000000	1000000000
Count:	192758	238362	80732	146139	39201	2207
Avg Freq:	697.59	2349.41	7300.64	38106.12	188165.50	1723067.42
(0, 3):	2355.34	6087.01	14133.42	41820.91	162365.18	1033792.12
(1, 2):	2217.11	6700.99	13895.22	40395.22	171603.17	1046385.28
(2, 1):	2223.30	6835.36	14281.67	42100.32	172929.41	1049067.41
(3, 0):	2397.17	6688.31	15609.27	44224.88	171740.53	1041373.26

Table 30: Seed length 13

Freqs:	0	5	10	50	100	500
Count:	119893	1296668	782215	1012545	169448	228469
Avg Freq:	0.00	2.90	7.71	21.54	70.75	231.81
(0, 5):	94.76	142.80	196.01	378.19	1050.82	2593.59
(1, 4):	88.25	137.32	196.19	365.75	1028.75	2575.29
(2, 3):	92.72	139.16	192.85	357.71	1051.37	2620.70
(3, 2):	93.05	139.42	187.77	364.28	1062.28	2669.95
(4, 1):	97.40	139.39	190.59	373.66	1074.14	2725.39
(5, 0):	92.20	143.02	198.27	391.13	1134.53	2930.94

Table 31: Seed length 15

Freqs:	1000	5000	10000	100000	1000000	1000000000
Count:	67449	147865	68003	110002	22731	1521
Avg Freq:	708.27	2513.11	7125.24	38015.44	173055.18	1191151.41
(0, 5):	6500.26	9933.95	11493.09	33278.47	123185.21	499685.63
(1, 4):	5368.47	9709.77	12161.65	33446.83	131134.11	509271.39
(2, 3):	5750.41	9122.51	12469.85	34774.53	134129.87	514628.16
(3, 2):	5787.67	9534.10	12678.14	34910.12	137436.07	518084.84
(4, 1):	5965.05	10939.14	13353.92	33978.20	142319.18	519248.13
(5, 0):	7576.56	11455.24	13135.02	34583.92	138267.35	513356.68

Table 32: Seed length 15

Freqs:	0	5	10	50	100	500
Count:	708601	2509784	107476	190386	67327	137416
Avg Freq:	0.00	1.35	7.64	24.33	72.14	239.50
(0, 8):	146.54	185.98	905.40	1727.28	2540.82	3525.44
(1, 7):	140.95	181.57	826.45	1576.29	2497.08	3581.61
(2, 6):	141.73	181.21	820.99	1522.09	2392.90	3517.53
(3, 5):	143.74	177.98	836.01	1550.87	2286.16	3400.05
(4, 4):	139.46	177.38	815.23	1568.57	2444.74	3401.16
(5, 3):	143.27	178.04	835.76	1571.80	2377.97	3628.36
(6, 2):	143.49	183.79	868.36	1620.57	2579.40	4011.43
(7, 1):	136.92	192.24	930.45	1831.74	2920.05	4070.44
(8, 0):	136.26	204.77	1092.63	2148.54	3045.26	4079.51

Table 33: Seed length 18

Freqs:	1000	5000	10000	100000	1000000	1000000000
Count:	45861	122776	49277	77158	10747	0
Avg Freq:	723.75	2557.84	6943.46	34782.67	211144.09	0.00
(0, 8):	5991.70	6760.99	8559.45	22857.12	99409.11	0.00
(1, 7):	6016.09	6928.29	8577.15	23204.92	102778.90	0.00
(2, 6):	6210.54	7092.48	8661.80	23809.86	104476.53	0.00
(3, 5):	6208.34	7406.93	8921.84	23769.40	107979.85	0.00
(4, 4):	6307.44	7736.48	9203.70	23643.40	111953.00	0.00
(5, 3):	6643.53	7744.20	9172.26	24029.35	111473.34	0.00
(6, 2):	6878.12	7708.36	9333.54	24259.54	110395.75	0.00
(7, 1):	6856.99	7757.74	9359.86	24430.85	108938.98	0.00
(8, 0):	7031.62	7812.77	9304.02	24638.03	107708.63	0.00

Table 34: Seed length 18

Freqs:	0	5	10	50	100	500
Count:	784079	2571704	80831	157567	57769	117097
Avg Freq:	0.00	1.18	7.71	24.65	72.27	235.91
(0, 10):	173.60	215.36	1055.95	1571.64	2173.87	3091.49
(1, 9):	169.24	204.65	1044.15	1574.85	2180.26	3130.75
(2, 8):	166.14	200.77	999.83	1568.93	2233.88	3167.77
(3, 7):	168.09	196.57	970.62	1549.92	2228.54	3193.01
(4, 6):	163.16	195.68	973.62	1571.02	2236.01	3252.70
(5, 5):	164.58	194.46	958.81	1592.64	2237.82	3312.36
(6, 4):	164.65	197.35	1009.97	1637.17	2406.66	3398.32
(7, 3):	162.80	201.84	1060.22	1739.83	2503.88	3447.56
(8, 2):	163.55	208.42	1140.96	1831.39	2557.90	3480.85
(9, 1):	162.41	223.73	1237.81	1855.60	2520.55	3498.96
(10, 0):	163.08	248.42	1301.65	1897.36	2534.19	3502.16

Table 35: Seed length 20

Freqs:	1000	5000	10000	100000	1000000	1000000000
Count:	40776	111922	37171	61976	5917	0
Avg Freq:	724.44	2578.90	7029.43	32161.38	204915.80	0.00
(0, 10):	4765.08	5345.54	7178.56	18479.25	82432.48	0.00
(1, 9):	4809.48	5413.55	7392.81	18658.86	82975.53	0.00
(2, 8):	4950.51	5468.12	7659.98	18811.26	84940.53	0.00
(3, 7):	5082.51	5624.04	7712.46	18808.93	87318.63	0.00
(4, 6):	5226.03	5804.72	7698.03	18856.53	89900.76	0.00
(5, 5):	5331.00	5916.56	7627.76	19003.34	90801.45	0.00
(6, 4):	5387.70	6004.87	7705.25	19110.16	90860.76	0.00
(7, 3):	5350.29	5996.10	7765.89	19356.38	89369.76	0.00
(8, 2):	5435.58	6007.84	7787.24	19644.72	87522.08	0.00
(9, 1):	5441.54	6029.58	7819.80	19807.91	85931.09	0.00
(10, 0):	5556.77	6043.95	7656.41	20133.16	85702.30	0.00

Table 36: Seed length 20

Freqs:	0	5	10	50	100	500
Count:	816832	2625940	72592	140757	52022	100360
Avg Freq:	0.00	1.15	7.70	24.66	72.72	236.19
(0, 12):	173.10	199.19	905.74	1343.88	1830.05	2675.92
(1, 11):	168.84	191.83	914.68	1337.32	1832.98	2696.23
(2, 10):	165.86	188.07	916.37	1352.53	1843.06	2720.94
(3, 9):	165.77	183.51	922.56	1365.61	1850.92	2762.44
(4, 8):	160.96	182.03	912.40	1410.90	1884.97	2820.11
(5, 7):	161.34	181.23	899.09	1424.98	1930.33	2881.49
(6, 6):	161.47	183.77	919.37	1447.82	2007.35	2945.90
(7, 5):	159.70	185.59	958.09	1484.91	2022.90	2962.68
(8, 4):	159.58	189.25	1022.16	1520.40	2046.71	2972.05
(9, 3):	159.35	196.29	1039.84	1536.76	2029.38	2971.38
(10, 2):	157.17	206.61	1051.87	1545.73	2045.62	2964.31
(11, 1):	155.64	218.98	1057.49	1554.12	2065.97	2977.41
(12, 0):	159.86	234.11	1074.41	1583.41	2096.32	3010.47

Table 37: Seed length 22

Freqs:	1000	5000	10000	100000	1000000	1000000000
Count:	36672	98988	32518	46541	3587	0
Avg Freq:	724.35	2563.87	7153.46	29390.73	181018.39	0.00
(0, 12):	3802.03	4268.56	5565.08	15264.58	62429.14	0.00
(1, 11):	3721.10	4315.08	5782.58	15227.76	62678.87	0.00
(2, 10):	3749.85	4371.60	5980.80	15209.59	64034.95	0.00
(3, 9):	3821.07	4435.22	6104.65	15112.36	65598.70	0.00
(4, 8):	3909.15	4482.56	6156.09	15060.06	67876.14	0.00
(5, 7):	4005.25	4486.61	6204.85	15042.99	69483.14	0.00
(6, 6):	4096.32	4528.20	6240.87	15019.70	70488.12	0.00
(7, 5):	4094.39	4559.67	6210.18	15162.97	70393.73	0.00
(8, 4):	4157.24	4596.58	6148.24	15358.89	69674.25	0.00
(9, 3):	4197.97	4618.91	6182.53	15536.38	68435.59	0.00
(10, 2):	4280.30	4617.48	6055.88	15849.58	66993.67	0.00
(11, 1):	4306.76	4611.93	5985.54	16047.22	65893.89	0.00
(12, 0):	4388.75	4625.24	5862.69	16356.85	65397.89	0.00

Table 38: Seed length 22

Freqs:	0	5	10	50	100	500
Count:	853607	2690186	63853	121559	45202	80411
Avg Freq:	0.00	1.13	7.69	24.45	72.22	237.72
(0, 15):	146.39	161.42	745.59	1117.90	1432.92	2175.54
(1, 14):	143.22	157.72	741.96	1113.32	1450.59	2191.55
(2, 13):	142.28	155.32	742.24	1113.50	1468.47	2211.74
(3, 12):	140.84	152.94	744.50	1113.80	1488.26	2231.81
(4, 11):	137.74	151.21	747.98	1136.73	1507.86	2264.47
(5, 10):	136.92	149.96	758.14	1156.87	1562.92	2278.56
(6, 9):	136.75	150.58	775.75	1175.56	1595.29	2293.14
(7, 8):	135.22	151.41	781.20	1191.83	1606.70	2292.46
(8, 7):	133.99	153.00	788.82	1207.48	1611.67	2320.52
(9, 6):	133.73	155.26	794.12	1205.93	1619.84	2336.90
(10, 5):	131.98	159.16	791.62	1202.04	1614.01	2339.32
(11, 4):	129.84	163.92	794.20	1202.98	1599.41	2342.90
(12, 3):	129.98	167.40	797.13	1208.91	1597.39	2346.62
(13, 2):	130.87	171.57	808.33	1206.11	1576.41	2343.47
(14, 1):	133.28	178.27	811.29	1216.69	1583.12	2350.08
(15, 0):	135.17	186.75	813.20	1236.31	1592.80	2373.69

Table 39: Seed length 25

Freqs:	1000	5000	10000	100000	1000000	1000000000
Count:	30319	84540	25750	29748	1634	0
Avg Freq:	720.23	2505.33	6938.21	26740.63	142364.21	0.00
(0, 15):	2909.74	2962.16	4382.68	11781.10	43068.11	0.00
(1, 14):	2881.89	3009.02	4435.95	11702.59	42475.91	0.00
(2, 13):	2898.62	3065.92	4492.42	11684.95	41972.42	0.00
(3, 12):	2938.71	3103.43	4507.86	11633.23	41922.99	0.00
(4, 11):	2942.50	3134.16	4548.19	11519.45	43714.50	0.00
(5, 10):	2933.02	3164.20	4530.14	11386.56	45431.60	0.00
(6, 9):	2950.71	3204.56	4504.49	11300.62	45858.71	0.00
(7, 8):	2963.75	3216.70	4504.29	11286.43	46313.43	0.00
(8, 7):	2976.91	3219.91	4467.72	11338.04	46442.20	0.00
(9, 6):	3021.66	3223.47	4460.72	11411.87	46276.79	0.00
(10, 5):	3054.27	3202.91	4472.30	11575.54	46084.15	0.00
(11, 4):	3059.21	3190.74	4472.23	11778.68	44664.61	0.00
(12, 3):	3095.02	3181.34	4460.39	11938.43	43212.25	0.00
(13, 2):	3125.38	3178.95	4446.57	12088.07	43382.08	0.00
(14, 1):	3151.53	3162.02	4414.87	12177.38	43944.70	0.00
(15, 0):	3183.09	3140.36	4422.06	12312.28	45009.59	0.00

Table 40: Seed length 25

15.2 Results of Frequency Ratio Test

Freqs:	0	1	5	10	50	100
Count:	3	21	349	1374	25497	79449
Avg Freq:	0.00	1.00	4.00	8.21	34.83	76.60
(0, 1):	0.00	11.10	5.72	3.41	2.55	2.28
(1, 0):	0.00	8.95	5.01	3.38	2.53	2.28

Table 41: Seed length 11

Freqs:	500	1000	5000	10000	100000	1000000000
Count:	396593	924687	1989959	261184	273495	74198
Avg Freq:	296.36	758.52	2063.03	6921.88	32482.25	311065.07
(0, 1):	1.85	1.38	1.24	1.16	1.06	1.01
(1, 0):	1.87	1.38	1.24	1.17	1.06	1.06

Table 42: Seed length 11

Freqs:	0	1	5	10	50	100
Count:	276	889	6963	15058	214606	155766
Avg Freq:	0.00	1.00	3.70	8.24	28.44	76.68
(0, 2):	0.00	16.04	10.40	7.33	5.07	4.00
(1, 1):	0.00	24.43	11.36	7.63	5.20	3.81
(2, 0):	0.00	18.23	9.17	7.12	5.17	3.95

Table 43: Seed length 12

Freqs:	500	1000	5000	10000	100000	1000000000
Count:	2045663	688840	555620	108938	179399	54791
Avg Freq:	276.09	693.54	2074.50	7095.21	36617.05	290419.39
(0, 2):	2.13	1.78	1.81	1.40	1.16	0.97
(1, 1):	2.13	1.74	1.78	1.50	1.20	1.04
(2, 0):	2.13	1.79	1.84	1.45	1.26	1.00

Table 44: Seed length 12

Freqs:	0	1	5	10	50	100
Count:	3802	10366	88371	98153	749457	1001262
Avg Freq:	0.00	1.00	3.42	7.87	33.10	73.31
(0, 3):	0.00	41.69	21.12	16.35	6.97	4.30
(1, 2):	0.00	56.91	19.05	14.05	6.99	4.28
(2, 1):	0.00	43.08	19.01	15.32	7.08	4.38
(3, 0):	0.00	58.21	18.63	16.44	7.05	4.34

Table 45: Seed length 13

Freqs:	500	1000	5000	10000	100000	1000000000
Count:	1375999	192758	238362	80732	146139	41408
Avg Freq:	206.49	697.59	2349.41	7300.64	38106.12	269974.05
(0, 3):	3.43	3.37	2.82	1.96	1.28	0.86
(1, 2):	3.36	3.16	3.03	1.93	1.15	0.90
(2, 1):	3.33	3.18	3.09	1.98	1.23	0.90
(3, 0):	3.48	3.44	3.06	2.19	1.42	0.91

Table 46: Seed length 13

Freqs:	0	1	5	10	50	100
Count:	119893	287608	1009060	782215	1012545	169448
Avg Freq:	0.00	1.00	3.45	7.71	21.54	70.75
(0, 5):	0.00	104.66	48.59	25.94	18.37	14.84
(1, 4):	0.00	104.44	46.24	25.88	17.93	14.48
(2, 3):	0.00	107.86	46.88	25.49	17.33	14.86
(3, 2):	0.00	106.20	47.02	24.87	17.67	15.08
(4, 1):	0.00	106.30	46.82	25.13	18.16	15.23
(5, 0):	0.00	105.99	48.37	26.19	18.84	16.19

Table 47: Seed length 15

Freqs:	500	1000	5000	10000	100000	1000000000
Count:	228469	67449	147865	68003	110002	24252
Avg Freq:	231.81	708.27	2513.11	7125.24	38015.44	236906.59
(0, 5):	11.64	9.26	4.70	1.64	1.01	0.70
(1, 4):	11.91	7.65	4.48	1.75	1.01	0.75
(2, 3):	12.03	8.19	4.31	1.78	1.06	0.77
(3, 2):	12.20	8.27	4.44	1.81	1.09	0.79
(4, 1):	12.54	8.52	5.05	1.92	1.07	0.82
(5, 0):	13.13	10.86	5.42	1.87	1.09	0.80

Table 48: Seed length 15

Freqs:	0	1	5	10	50	100
Count:	708601	1996228	513556	107476	190386	67327
Avg Freq:	0.00	1.00	2.70	7.64	24.33	72.14
(0, 8):	0.00	142.18	132.59	119.70	79.80	36.34
(1, 7):	0.00	140.53	126.94	109.25	70.90	35.57
(2, 6):	0.00	140.14	126.18	108.76	69.43	34.02
(3, 5):	0.00	137.29	123.64	110.86	70.14	32.67
(4, 4):	0.00	137.10	122.66	108.22	71.06	34.86
(5, 3):	0.00	137.23	123.01	110.78	71.11	33.97
(6, 2):	0.00	141.04	129.65	115.66	74.10	36.66
(7, 1):	0.00	146.91	136.27	124.14	82.55	41.80
(8, 0):	0.00	155.12	146.70	144.72	99.73	43.69

Table 49: Seed length 18

Freqs:	500	1000	5000	10000	100000	1000000000
Count:	137416	45861	122776	49277	77158	10747
Avg Freq:	239.50	723.75	2557.84	6943.46	34782.67	211144.09
(0, 8):	16.79	8.51	3.22	1.25	0.72	0.54
(1, 7):	16.94	8.53	3.30	1.25	0.74	0.56
(2, 6):	16.58	8.80	3.38	1.27	0.76	0.57
(3, 5):	15.81	8.79	3.50	1.31	0.76	0.60
(4, 4):	15.88	8.93	3.68	1.34	0.75	0.63
(5, 3):	16.88	9.46	3.68	1.35	0.76	0.62
(6, 2):	19.05	9.80	3.69	1.37	0.77	0.62
(7, 1):	19.50	9.77	3.73	1.38	0.78	0.61
(8, 0):	19.69	10.01	3.77	1.37	0.78	0.60

Table 50: Seed length 18

Freqs:	0	1	5	10	50	100
Count:	784079	2314973	256731	80831	157567	57769
Avg Freq:	0.00	1.00	2.84	7.71	24.65	72.27
(0, 10):	0.00	164.21	244.45	140.79	72.51	30.96
(1, 9):	0.00	157.75	226.81	138.92	72.59	31.10
(2, 8):	0.00	157.05	215.55	132.67	72.40	31.83
(3, 7):	0.00	153.75	211.97	128.71	70.93	31.74
(4, 6):	0.00	152.96	211.41	128.88	71.63	31.87
(5, 5):	0.00	151.73	210.92	126.64	72.77	31.83
(6, 4):	0.00	154.96	211.52	134.00	74.82	34.06
(7, 3):	0.00	158.44	215.87	141.06	79.90	35.48
(8, 2):	0.00	162.43	225.58	151.04	85.22	36.43
(9, 1):	0.00	171.10	252.11	164.35	87.25	35.92
(10, 0):	0.00	188.48	287.06	173.37	89.71	35.99

Table 51: Seed length 20

Freqs:	500	1000	5000	10000	100000	1000000000
Count:	117097	40776	111922	37171	61976	5917
Avg Freq:	235.91	724.44	2578.90	7029.43	32161.38	204915.80
(0, 10):	14.69	6.78	2.58	1.05	0.62	0.47
(1, 9):	14.86	6.87	2.59	1.08	0.63	0.47
(2, 8):	15.05	7.07	2.61	1.12	0.64	0.48
(3, 7):	15.17	7.26	2.69	1.13	0.64	0.49
(4, 6):	15.42	7.47	2.79	1.12	0.63	0.51
(5, 5):	15.59	7.66	2.85	1.11	0.64	0.51
(6, 4):	16.14	7.74	2.90	1.12	0.64	0.52
(7, 3):	16.39	7.70	2.89	1.13	0.65	0.51
(8, 2):	16.58	7.81	2.89	1.14	0.66	0.49
(9, 1):	16.72	7.81	2.91	1.14	0.66	0.48
(10, 0):	16.72	7.95	2.94	1.12	0.67	0.48

Table 52: Seed length 20

Freqs:	0	1	5	10	50	100
Count:	816832	2416456	209484	72592	140757	52022
Avg Freq:	0.00	1.00	2.91	7.70	24.66	72.72
(0, 12):	0.00	160.84	235.37	120.75	62.33	26.12
(1, 11):	0.00	154.26	228.54	122.13	62.00	26.24
(2, 10):	0.00	151.18	222.88	122.63	62.73	26.35
(3, 9):	0.00	146.82	220.31	123.34	63.50	26.40
(4, 8):	0.00	145.83	217.16	122.12	65.60	26.92
(5, 7):	0.00	145.44	215.59	119.89	66.17	27.54
(6, 6):	0.00	147.33	219.91	122.49	67.25	28.65
(7, 5):	0.00	148.44	223.33	127.27	69.23	28.86
(8, 4):	0.00	150.48	231.65	135.90	71.12	29.17
(9, 3):	0.00	155.93	240.53	138.10	72.08	28.87
(10, 2):	0.00	164.93	250.36	140.01	72.24	29.09
(11, 1):	0.00	176.55	259.85	140.91	72.69	29.38
(12, 0):	0.00	190.02	274.22	143.04	74.33	29.83

Table 53: Seed length 22

Freqs:	500	1000	5000	10000	100000	1000000000
Count:	100360	36672	98988	32518	46541	3587
Avg Freq:	236.19	724.35	2563.87	7153.46	29390.73	181018.39
(0, 12):	12.65	5.44	2.09	0.80	0.55	0.39
(1, 11):	12.71	5.34	2.10	0.83	0.55	0.39
(2, 10):	12.89	5.37	2.12	0.86	0.55	0.40
(3, 9):	13.09	5.48	2.16	0.88	0.55	0.41
(4, 8):	13.32	5.59	2.17	0.89	0.55	0.42
(5, 7):	13.68	5.73	2.18	0.90	0.55	0.43
(6, 6):	14.01	5.87	2.20	0.90	0.55	0.44
(7, 5):	14.10	5.87	2.22	0.90	0.55	0.44
(8, 4):	14.16	5.95	2.24	0.89	0.56	0.43
(9, 3):	14.16	6.01	2.25	0.89	0.56	0.43
(10, 2):	14.17	6.12	2.25	0.87	0.57	0.42
(11, 1):	14.23	6.15	2.25	0.86	0.58	0.41
(12, 0):	14.36	6.26	2.28	0.84	0.59	0.40

Table 54: Seed length 22

Freqs:	0	1	5	10	50	100
Count:	853607	2504219	185967	63853	121559	45202
Avg Freq:	0.00	1.00	2.91	7.69	24.45	72.22
(0, 15):	0.00	136.21	183.55	98.93	52.06	20.38
(1, 14):	0.00	132.46	182.14	98.57	51.70	20.67
(2, 13):	0.00	130.11	181.36	98.50	51.95	20.93
(3, 12):	0.00	127.49	182.09	98.81	52.00	21.20
(4, 11):	0.00	125.69	181.43	99.24	53.29	21.50
(5, 10):	0.00	124.60	179.77	100.66	54.36	22.26
(6, 9):	0.00	125.12	179.72	103.08	55.19	22.74
(7, 8):	0.00	125.44	182.40	103.78	55.97	22.88
(8, 7):	0.00	126.19	187.11	105.28	56.91	22.96
(9, 6):	0.00	128.32	188.02	105.90	56.95	23.09
(10, 5):	0.00	132.06	190.64	105.68	56.58	23.00
(11, 4):	0.00	136.63	193.59	106.01	56.71	22.79
(12, 3):	0.00	140.36	194.40	106.25	56.91	22.75
(13, 2):	0.00	144.45	197.25	107.47	56.61	22.44
(14, 1):	0.00	151.15	199.61	107.78	57.16	22.58
(15, 0):	0.00	159.28	204.48	108.34	58.01	22.72

Table 55: Seed length 25

Freqs:	500	1000	5000	10000	100000	1000000000
Count:	80411	30319	84540	25750	29748	1634
Avg Freq:	237.72	720.23	2505.33	6938.21	26740.63	142364.21
(0, 15):	10.39	4.15	1.45	0.64	0.45	0.30
(1, 14):	10.48	4.12	1.46	0.65	0.45	0.30
(2, 13):	10.61	4.13	1.49	0.66	0.45	0.30
(3, 12):	10.72	4.20	1.51	0.67	0.45	0.30
(4, 11):	10.88	4.21	1.52	0.67	0.44	0.31
(5, 10):	10.98	4.20	1.53	0.67	0.44	0.33
(6, 9):	11.09	4.23	1.55	0.67	0.44	0.33
(7, 8):	11.10	4.25	1.55	0.67	0.44	0.33
(8, 7):	11.24	4.26	1.56	0.66	0.44	0.34
(9, 6):	11.32	4.33	1.56	0.66	0.44	0.33
(10, 5):	11.30	4.38	1.55	0.66	0.45	0.33
(11, 4):	11.30	4.39	1.55	0.66	0.46	0.32
(12, 3):	11.30	4.44	1.55	0.66	0.46	0.31
(13, 2):	11.25	4.47	1.55	0.65	0.47	0.31
(14, 1):	11.25	4.51	1.54	0.65	0.47	0.31
(15, 0):	11.34	4.55	1.54	0.65	0.48	0.32

Table 56: Seed length 25

15.3 Results of Frequency Comparison Test

(left, right)	% of correct comparisons
(0, 1):	0.885686
(1, 0):	0.885490

Table 57: Seed length 11

(left, right)	% of correct comparisons
(0, 2):	0.838487
(1, 1):	0.834117
(2, 0):	0.838308

Table 58: Seed length 12

(left, right)	% of correct comparisons
(0, 3):	0.802248
(1, 2):	0.799497
(2, 1):	0.799555
(3, 0):	0.801996

Table 59: Seed length 13

(left, right)	% of correct comparisons
(0, 5):	0.734238
(1, 4):	0.733823
(2, 3):	0.734016
(3, 2):	0.734292
(4, 1):	0.733973
(5, 0):	0.735255

Table 60: Seed length 15

(left, right)	% of correct comparisons
(0, 8):	0.644817
(1, 7):	0.648471
(2, 6):	0.647196
(3, 5):	0.647320
(4, 4):	0.649201
(5, 3):	0.649550
(6, 2):	0.649888
(7, 1):	0.649108
(8, 0):	0.649469

Table 61: Seed length 18

(left, right)	% of correct comparisons
(0, 10):	0.627535
(1, 9):	0.628668
(2, 8):	0.628174
(3, 7):	0.628991
(4, 6):	0.629764
(5, 5):	0.629845
(6, 4):	0.630774
(7, 3):	0.630830
(8, 2):	0.631497
(9, 1):	0.631189
(10, 0):	0.630728

Table 62: Seed length 20

(left, right)	% of correct comparisons
(0, 12):	0.616800
(1, 11):	0.617805
(2, 10):	0.617108
(3, 9):	0.617824
(4, 8):	0.618703
(5, 7):	0.618959
(6, 6):	0.619920
(7, 5):	0.619714
(8, 4):	0.620441
(9, 3):	0.621070
(10, 2):	0.621115
(11, 1):	0.619363
(12, 0):	0.620226

Table 63: Seed length 22

(left, right)	% of correct comparisons
(0, 15):	0.605822
(1, 14):	0.606799
(2, 13):	0.605634
(3, 12):	0.606399
(4, 11):	0.606932
(5, 10):	0.607885
(6, 9):	0.608510
(7, 8):	0.609034
(8, 7):	0.609307
(9, 6):	0.610289
(10, 5):	0.610546
(11, 4):	0.609536
(12, 3):	0.610421
(13, 2):	0.609668
(14, 1):	0.608081
(15, 0):	0.606699

Table 64: Seed length 25

16 Appendix B

16.1 Seed Selection Algorithm Speedups

Threads	Runtime (s)	Speedup
1	0.647	1.00
2	1.078	0.60
4	0.751	0.86
8	0.132	5.26
16	0.321	2.02
32	0.279	2.32
64	0.396	1.63
80	0.648	1.00

Table 65: Speedup of Uniform Selector (4M read set)

Threads	Runtime (s)	Speedup
1	73.49	1.00
2	39.90	1.84
4	21.14	3.48
8	10.42	7.05
16	5.34	13.76
32	3.34	22.00
64	1.85	39.73
80	1.70	43.23

Table 66: Speedup of Centered Bidirectional Uniform Selector (4M read set)

Threads	Runtime (s)	Speedup
1	621.69	1.00
2	341.45	1.82
4	169.80	3.66
8	82.08	7.57
16	40.99	15.17
32	20.34	30.56
64	10.91	56.98
80	9.27	67.06

Table 67: Speedup of Exact Centered Bidirectional Uniform Selector (4M read set)

Threads	Runtime (s)	Speedup
1	80.85	1.00
2	40.61	1.99
4	21.36	3.79
8	10.94	7.39
16	6.08	13.30
32	4.92	16.43
64	1.59	50.85
80	1.33	60.79

Table 68: Speedup of Hobbes Selector (4M read set)

Threads	Runtime (s)	Speedup
1	107.69	1.00
2	54.21	1.98
4	39.09	2.75
8	13.17	8.18
16	6.72	16.02
32	3.44	31.31
64	2.09	51.53
80	1.78	60.50

Table 69: Speedup of Bidirectional Hobbes Selector (4M read set)

Threads	Runtime (s)	Speedup
1	190.36	1.00
2	94.70	2.01
4	53.85	3.54
8	23.37	8.15
16	11.77	16.17
32	6.22	30.60
64	3.85	49.44
80	3.17	60.05

Table 70: Speedup of Centered Bidirectional Hobbes Selector (4M read set)

Threads	Runtime (s)	Speedup
1	189.35	1.00
2	95.89	1.97
4	55.45	3.41
8	24.48	7.73
16	12.00	15.78
32	7.68	24.66
64	3.61	52.45
80	2.95	64.19

Table 71: Speedup of Exact Bidirectional Hobbes Selector (4M read set)

Threads	Runtime (s)	Speedup
1	605.79	1.00
2	358.56	1.69
4	169.77	3.57
8	81.20	7.46
16	40.60	14.92
32	20.63	29.36
64	11.23	53.84
80	9.30	65.14

Table 72: Speedup of Exact Centered Bidirectional Hobbes Selector (4M read set)